



Direct PixelNav: red = fixed initial goal mask; yellow = model-predicted goal on STOP frame
BGR to RGB display matches preprocessing. Visual markers are not depth/ground truth.


direct_goal-goal-4-002 initial RGB + goal mask



Actual 224x224 policy goal input



Model STOP; residual 7.46 m



direct_goal-goal-4-003 initial RGB + goal mask




Actual 224x224 policy goal input




Model STOP; residual 5.66 m



direct_goal-goal-4-004 initial RGB + goal mask




Actual 224x224 policy goal input



Model STOP; residual 9.35 m




direct_goal-goal-4-005 initial RGB + goal mask




Actual 224x224 policy goal input




Model STOP; residual 8.79 m



direct_goal-goal-4-006 initial RGB + goal mask



Actual 224x224 policy goal input



Model STOP; residual 9.47 m

